

(Don't) Plan To Escape the Maze

Konstantin Pakulev, Mohammed Deifallah,
Hekmat Taherinejad, Pavel Krasnik






Project objectives


PS4 (*Escape the Maze*) is taken as a basis for our work

- Implement different planning algorithms
 - Value Iteration
 - Markov Decision Process
 - Monte-Carlo Tree Search
- Implement different policies for pursuers
- Explore their performance under
 - Different **number** of pursuers
 - Various **operating logic** of pursuers

Easy mode



Hard mode



Value Iteration

Method formulation:


$$G_k^*(x_k) = \min_{u_k} \{l(x_k, u_k) + G_{k+1}^*(x_{k+1})\}$$

where $x_{k+1} = f(x_k, u_k)$


*optimal
cost-to-go*

$$u^* = \arg \min_{u \in U(x)} \{l(x, u) + G^*(f(x, u))\}$$

*recover the
optimal plan*




$G^*(x)$



VI (gets caught) payout


Improved VI

To avoid getting caught **the cost of the pursuers** should be taken into account




VI + pursuer cost
(escapes) payout

Improved Value Iteration



Pursuers cost smoothed by a gaussian kernel




$G^*(f(x,u)) + P(f(x,u))$


The recovery of the “optimal” plan then becomes:

$$u^* = \arg \min_{u \in U(x)} \left\{ G^*(f(x,u)) + P(f(x,u)) \right\}$$

Failure examples



VI + pursuer cost (gets caught) payout



VI + pursuer cost (gets stuck) payout

Markov Decision Process

Method formulation:

Bellman optimality equation:

$$v_*(s) = \max_{u \in U(s)} q_*(s, u) = \max_{u \in U(s)} \sum_{s', r} p(s', r | s, u) [r + \gamma v_*(s')]$$

After which the policy can be recovered by


$$u^* = \arg \max_{u \in U(s)} q_*(s, u)$$

For our problem we simplified the equation to:


$$v_*(s) = r + \gamma \max_{u \in U(s)} \sum_{s'} p(s' | s, u) v_*(s')$$

And the policy is changed to:

$$u^* = \arg \max_{u \in U(s)} -q_*(s, u) + R(s, u)$$




$R(s, u)$




Solving the Bellman equation
doesn't influence **pursuers**


Markov Decision Process. Pursuers




Locations of **pursuers**




$q_*(s, u)$



$-q_*(s, u) + R(s, u)$




MDP (**escapes**)
layout



MDP tricks the
pursuer to escape

MDP (**gets stuck**)
playout:



Monte-Carlo Tree Search

Reward function from the PS4:

$$R_T(x_e, x_p) = \begin{cases} 0 & \text{if } x_e = x_p \\ 100 + \sum_t^T \frac{0.1}{d(x_e(t) - x_{goal})} - T & \text{if } x_e = x_{goal} \\ \sum_t^T \frac{0.1}{d(x_e(t) - x_{goal})} - T & \text{otherwise} \end{cases}$$


Simulations with **high number of iterations** in our environment have **high chance to fail** providing **no information**

Solutions:


- **Relax** the failure reward
- Make **shorter simulation**

In our case the reward is modified to:

$$R_T(x_e, x_p) = \begin{cases} -100 + \sum_t^T \frac{0.1}{d(x_e(t) - x_{goal})} & \text{if } x_e = x_p \\ 100 + \sum_t^T \frac{0.1}{d(x_e(t) - x_{goal})} - T & \text{if } x_e = x_{goal} \\ \sum_t^T \frac{0.1}{d(x_e(t) - x_{goal})} - T & \text{otherwise} \end{cases}$$




MCTS (escapes)
payout



MCTS (escapes) payout
with less simulation
iterations

MCTS (gets stuck)
payout:



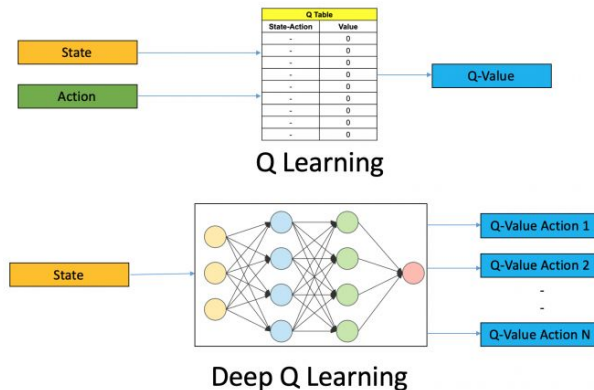
In the Q-learning algorithm, we compute the Q-table which contains the Q-values of any state-action pair using the Q-value iteration. In deep Q-learning, we use a neural network to approximate the Q-value function. The state is given as the input and the Q-value of all possible actions is generated as the output.

$$Q(s, a) = r(s, a) + \gamma \max_a Q(s', a)$$

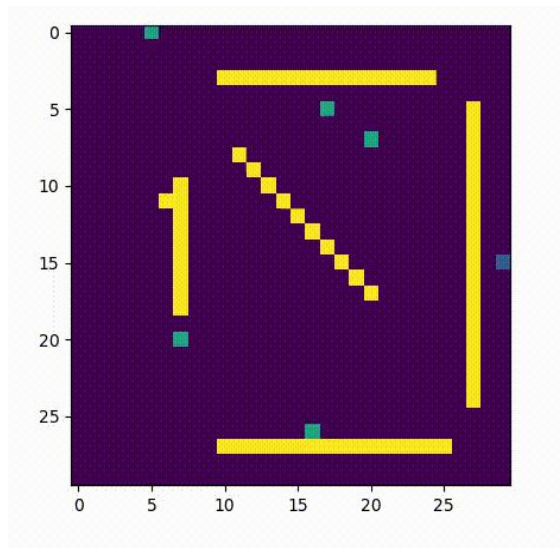
$$Q(s, a) \rightarrow \gamma Q(s', a) + \gamma^2 Q(s'', a) \dots \dots \gamma^n Q(s'' \dots n, a)$$

$$Q(S_t, A_t) = (1 - \alpha) Q(S_t, A_t) + \alpha * (R_t + \lambda * \max_a Q(S_{t+1}, a))$$

$$L_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \rho(\cdot)} [(y_i - Q(s, a; \theta_i))^2] \text{ where } y_i = r + \gamma \max_{a'} Q(s', a'; \theta_{i-1})$$




Variants of pursuers




Pursuers with **heuristic**
(manhattan distance)

$$d(x, y) = \sum_{i=0}^n |x_i - y_i|$$



Pursuers with **heuristic**
(euclidean distance)

$$d(x, y) = \sqrt{\sum_{i=0}^n (x_i - y_i)^2}$$



Pursuers with **Probabilistic**
Roadmap

Sample-based. Agents
act as a whole




Probabilistic Roadmaps (Pursuer)

Idea:

1. The agent is the set of pursuers as whole.
2. Sample C-space.
3. Get approximate result via graph search.
4. At proximity use accurate algorithm (A*).

Disadvantages:

- Curse of dimensionality: we need $\sim 5^6$ samples for 3 pursuers
- Too much for graph search
- Long initialization
- Random process






Probabilistic Roadmaps (Pursuer)

Advantages:


1. Simplification of space.
2. Easy to configure sophisticated constraints
3. Due to randomness result looks natural

Conclusion:


- Very niche algorithm
- Using lattice instead of uniform looks promising




(Not) Escaping the maze



MVI (gets caught)
vs
Pursuers with heuristic
(euclidean distance)



MDP (gets caught)
vs
Pursuers with heuristic
(euclidean distance)



MCTS (gets caught)
vs
Pursuers with heuristic
(euclidean distance)

Conclusions

- In the presence of **pursuers** with a just a little **more complex logic** the **difficulty increases dramatically**
- For both **MDP** and **MCTS** the choice of the **reward function** significantly changes the **behaviour** of the **escaper**
- **MDP** is not very suitable for **long term minimization/maximization** planning
- **MCTS** is noticeably faster than **MDP**